



deeplearning.ai

Case Studies

Why look at
case studies?

Outline

Classic networks:

- LeNet-5 ←
- AlexNet ←
- VGG ←

ResNet (152)

Inception




deeplearning.ai

Case Studies

Classic networks

LeNet - 5



60k parameters.

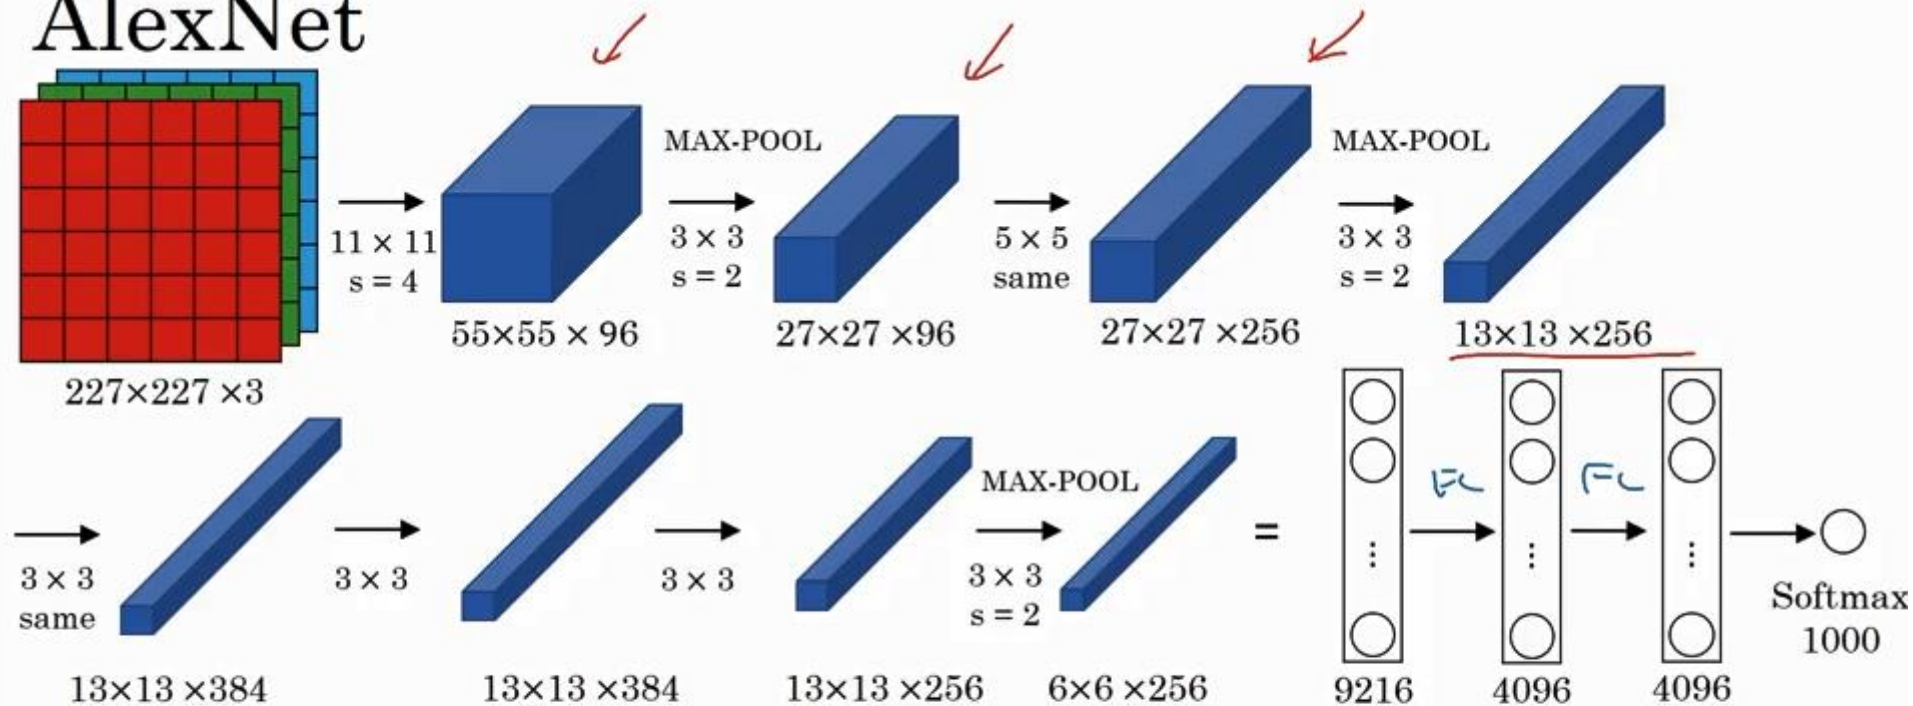
$n_H, n_W \downarrow$ $n_C \uparrow$

conv pool conv pool fc fc output

Advanced: sigmoid/tanh ReLU

II III

AlexNet



- Similar to LeNet, but much bigger.

- ReLU

- Multiple GPUs.

- Local Response Normalization (LRN)



~60M parameters

[Krizhevsky et al., 2012. ImageNet classification with deep convolutional neural networks]



Andrew Ng

VGG - 16




VGG-19

CONV = 3x3 filter, s = 1, same

MAX-POOL = 2x2, s = 2



224x224 x 3



$n_h, n_w \downarrow$ $n_c \uparrow$

~138M

[Simonyan & Zisserman 2015. Very deep convolutional networks for large-scale image recognition]

Andrew Ng




deeplearning.ai


Case Studies

Residual Networks (ResNets)

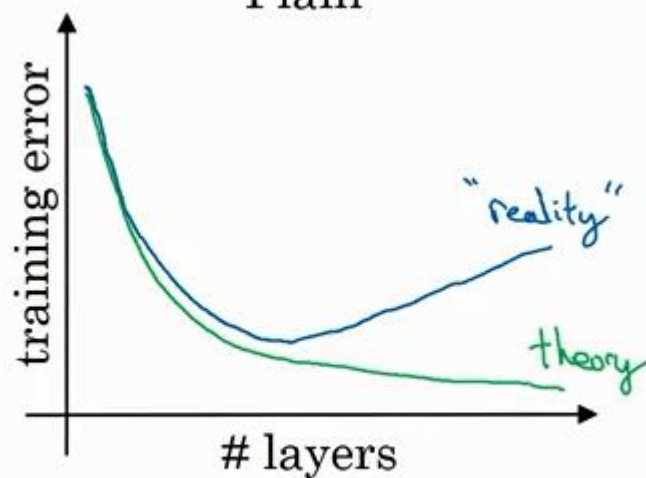
Residual block




Residual Network



Plain



ResNet



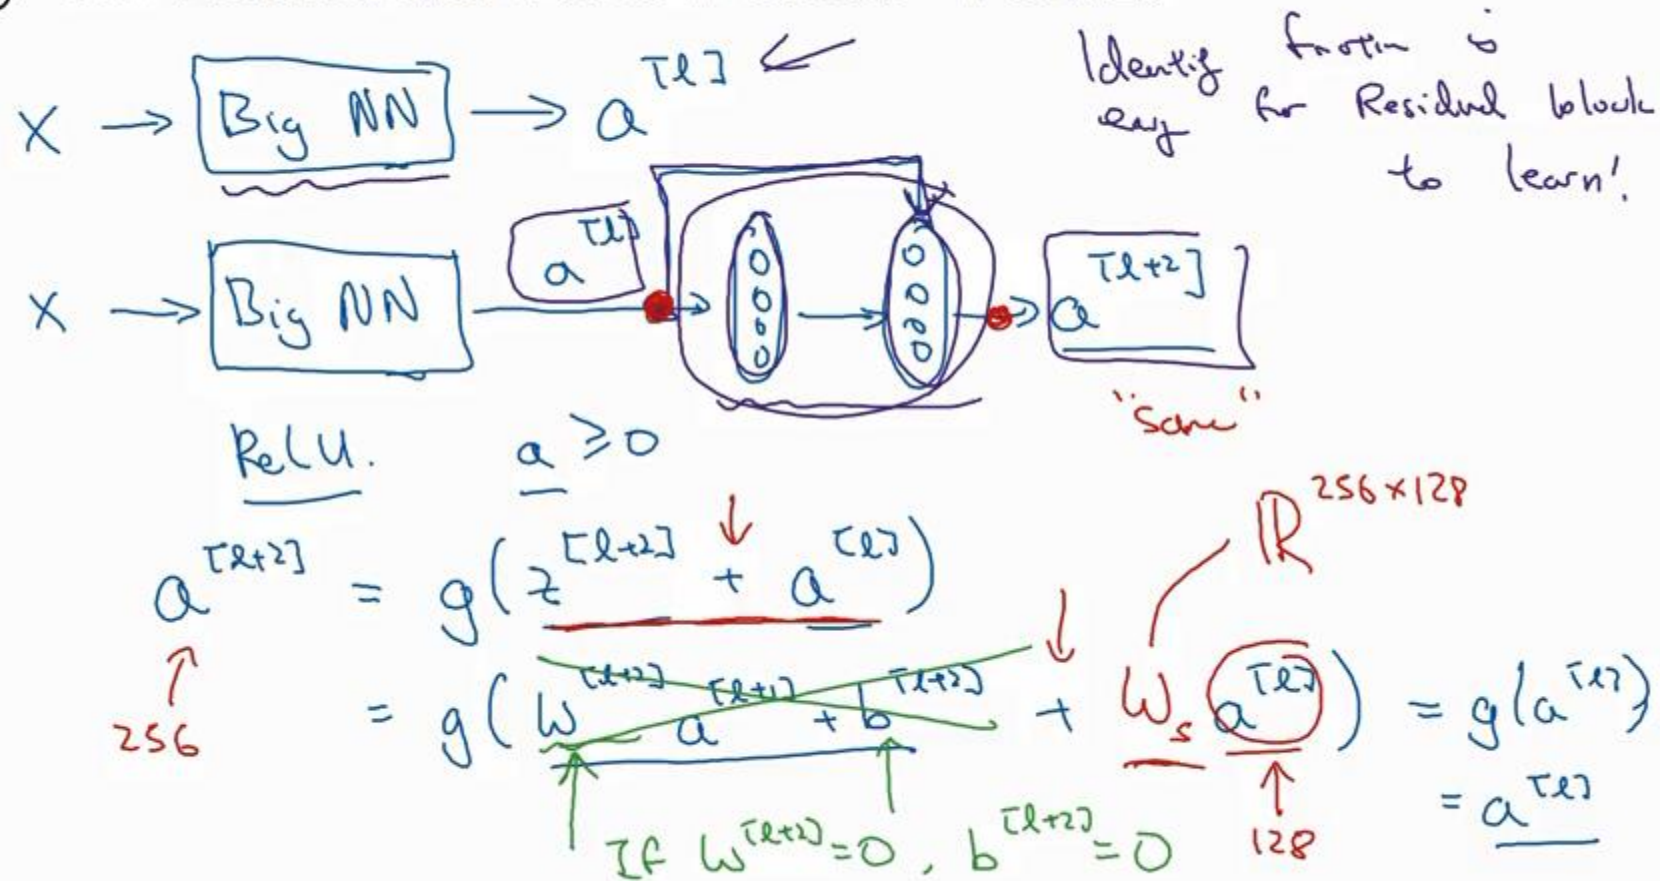


deeplearning.ai


Case Studies

Why ResNets work

Why do residual networks work?



ResNet





deeplearning.ai


Case Studies

Network in Network
and 1×1 convolutions

Why does a 1×1 convolution do?

1	2	3	6	5	8
3	5	5	1	3	4
2	1	3	4	9	3
4	7	8	5	7	9
1	5	3	7	4	8
5	4	9	8	3	5

$6 \times 6 \times 1$



$6 \times 6 \times 32$


*

2

=



32 \rightarrow # filters $n_c^{[L+1]}$



$1 \times 1 \times 32$

=


ReLU
Network in Network

2	4	6	...		

$6 \times 6 \times \# \text{ filters}$

Andrew Ng

Using 1×1 convolutions



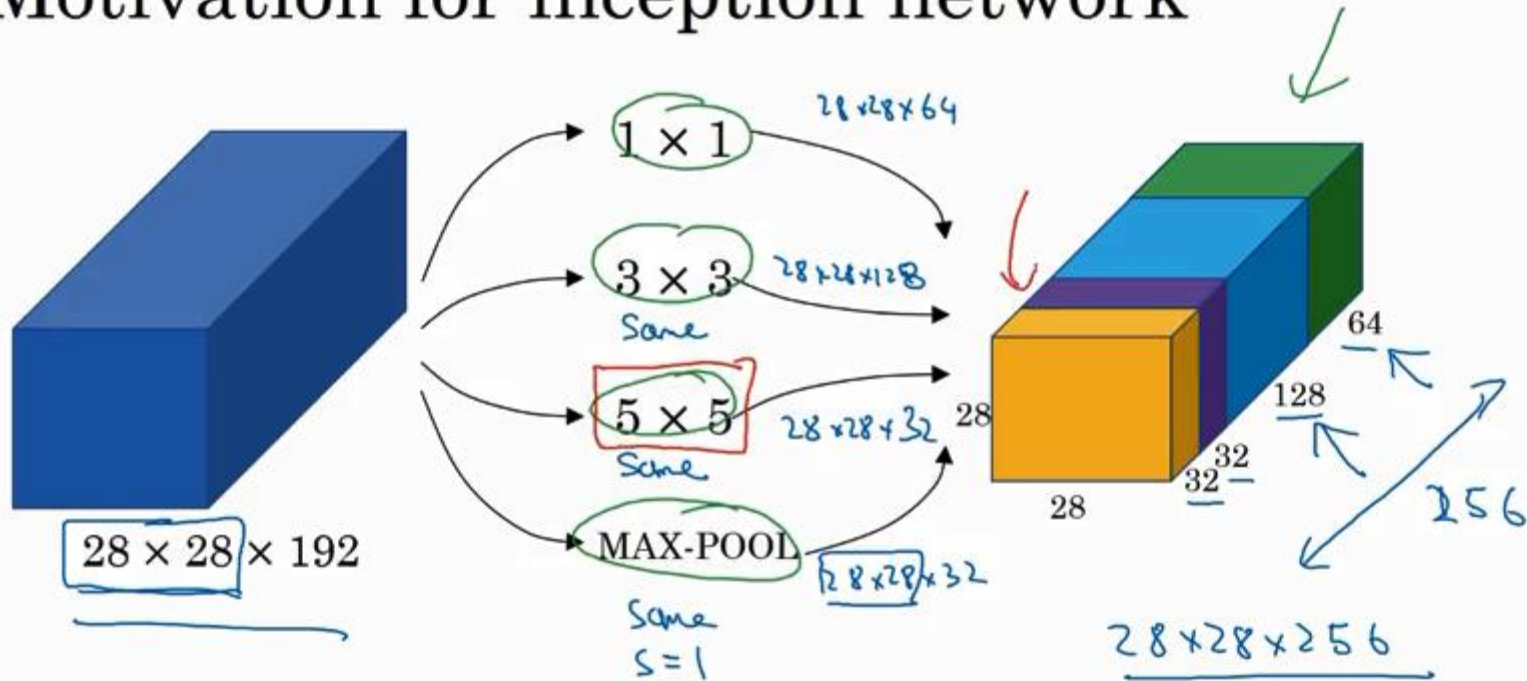


deeplearning.ai

Case Studies

Inception network motivation


Motivation for inception network




[Szegedy et al. 2014. Going deeper with convolutions]

Andrew Ng

The problem of computational cost



Using 1x1 convolution



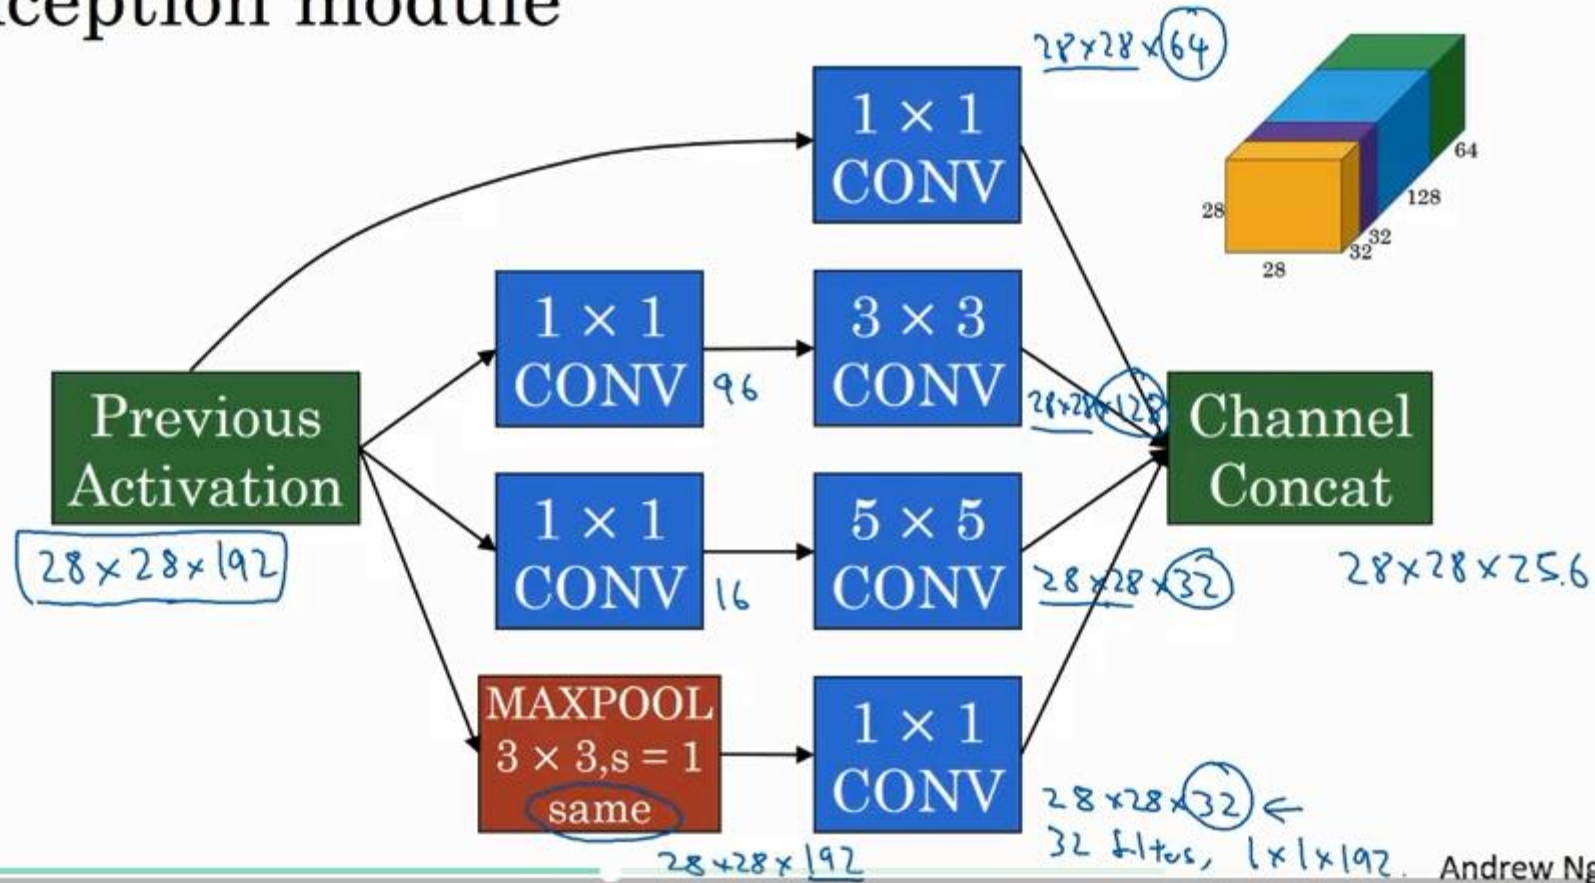


deeplearning.ai

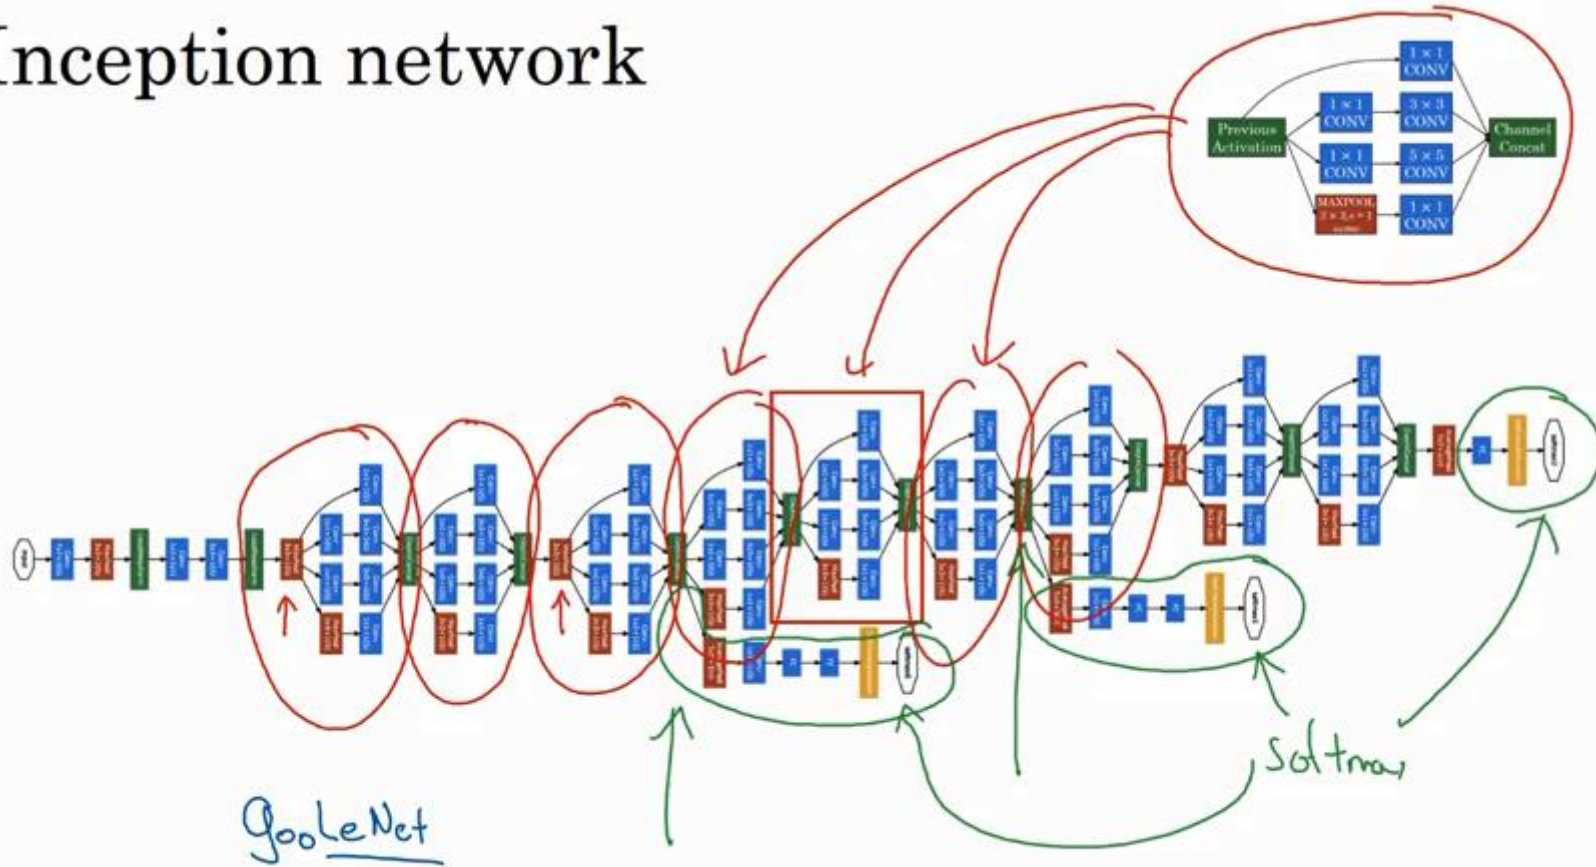
Case Studies

Inception network

Inception module



Inception network



[Szegedy et al., 2014, Going Deeper with Convolutions]

Andrew Ng



<http://knowyourmeme.com/memes/we-need-to-go-deeper> ←

Andrew Ng



deeplearning.ai

Practical advice for
using ConvNets

Using open-source
implementations






deeplearning.ai

Practical advice for
using ConvNets

Transfer Learning

Transfer Learning

 Tigger
  Misty
  Neither



Andrew Ng



deeplearning.ai

Practical advice for
using ConvNets

Data augmentation

Common augmentation method

Mirroring



Random Cropping



Rotation

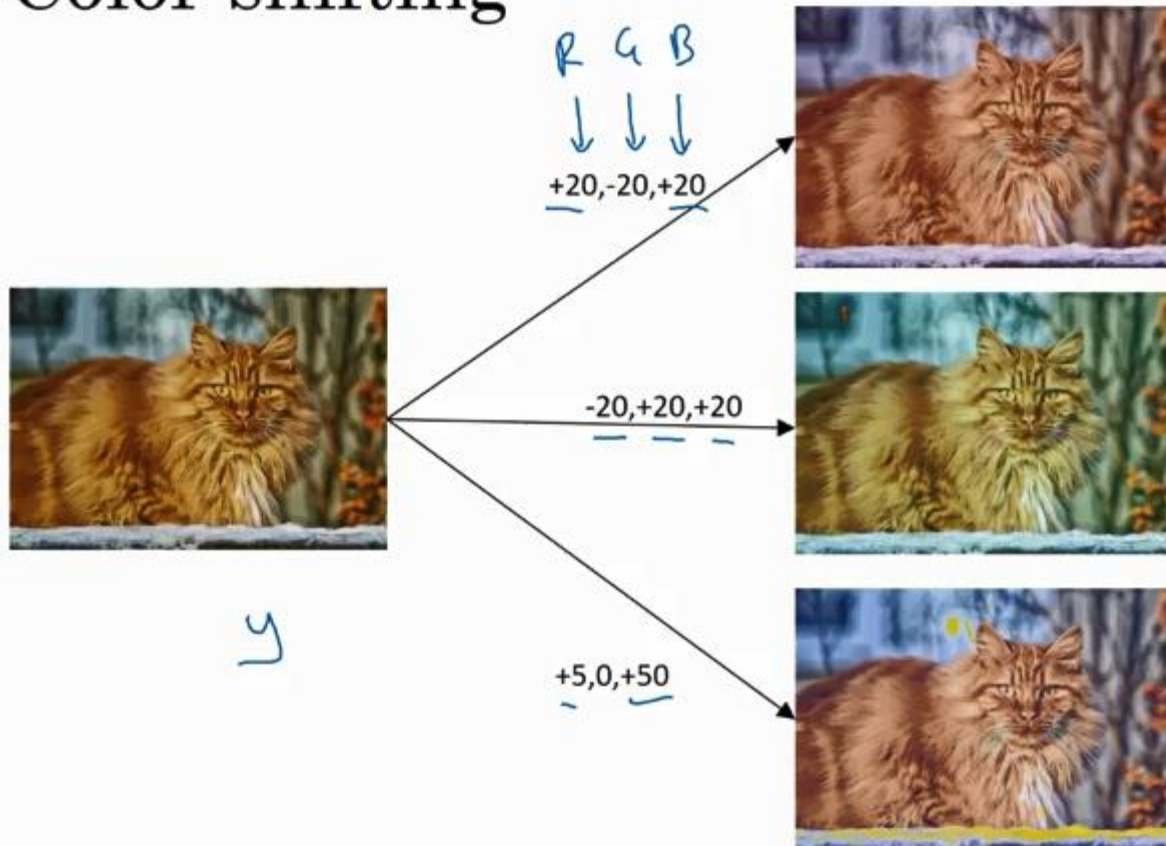
Shearing

Local warping

...



Color shifting



Advanced:

PCA

ml-class.org


[AlexNet paper

["PCA color augmentation."

R B

G

Implementing distortions during training



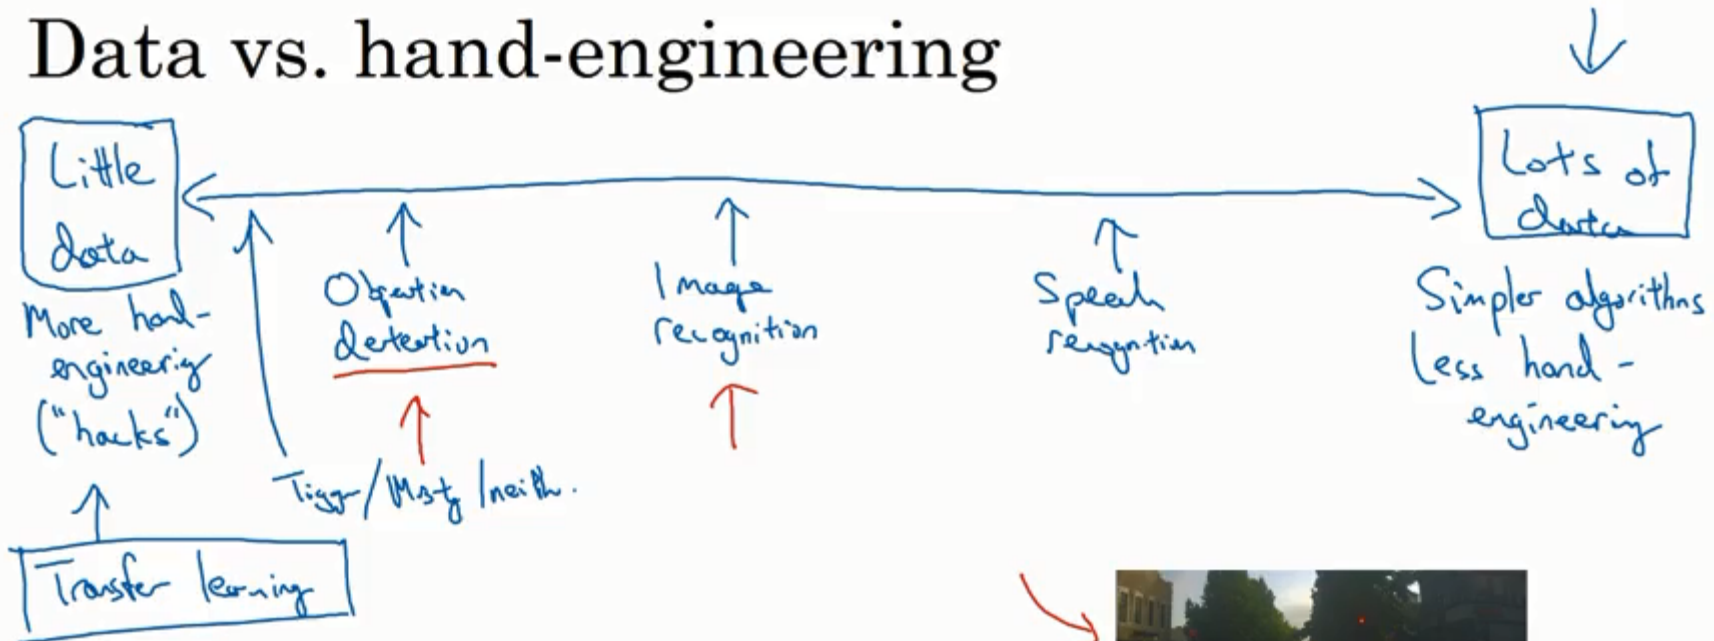


deeplearning.ai

Practical advice for
using ConvNets

The state of
computer vision

Data vs. hand-engineering



Two sources of knowledge

- • Labeled data (x, y)
- • Hand engineered features/network architecture/other components



Andrew Ng

Tips for doing well on benchmarks/winning competitions

网易云课堂

Ensembling

3-15 networks

$\rightarrow \hat{y}$

- Train several networks independently and average their outputs

Multi-crop at test time

- Run classifier on multiple versions of test images and average results

10-crop



1

+



4

+



1

+



4

Andrew Ng

Use open source code

- Use architectures of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset